

Conference Paper Title*

*Note: Sub-titles are not captured for <https://ieeexplore.ieee.org> and should not be used

1st Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

2nd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

3rd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

4th Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

5th Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

6th Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

Abstract—This document is a model and instructions for **LaTeX**. This and the **IEEEtran.cls** file define the components of your paper [title, text, heads, etc.]. ***CRITICAL: Do Not Use Symbols, Special Characters, Footnotes, or Math in Paper Title or Abstract.**

Index Terms—component, formatting, style, styling, insert.

I. INTRODUCTION

The rapid advancement of autonomous driving systems relies heavily on the accuracy and robustness of perception modules, particularly object detection [1], [2]. Historically, industry standards have favored closed-set object detection models, such as the widely adopted You Only Look Once architectures [3], [4]. These models are trained on large, meticulously annotated datasets to achieve real-time inference and high precision for predefined categories. However, their operational capacity is strictly bounded by the classes explicitly registered during the training phase.

Real-world driving environments are highly dynamic and unpredictable, continuously presenting edge cases and novel objects that fall outside the domain of standard training distributions. To benchmark perception models, the research community relies on various large-scale driving datasets, which can be grouped by their specific environmental focuses. For instance, Cityscapes [?] focuses on the high-resolution semantic understanding of complex urban street scenes across different cities. The India Driving Dataset [?] provides unstructured and highly unconstrained traffic environments typical of developing regions. To account for weather and lighting variations, ACDC [?] specifically targets adverse conditions such as fog, nighttime, rain, and snow. Meanwhile, BDD100K [5] serves as one of the most comprehensive datasets, offering diverse driving scenarios, weather conditions, and times of day to support heterogeneous multitask learning. Despite the extensive coverage of these datasets, when traditional closed-set models encounter out-of-distribution objects not present

in their specific training distributions, they inherently fail to detect them or force a misclassification into a known category [6], [7]. In the context of autonomous navigation, this fundamental limitation poses severe safety hazards, as the vehicle planning systems remain blind to critical real-world obstacles and corner cases [8], [9].

To mitigate the vulnerabilities of closed-set paradigms, the field is increasingly shifting towards Vision-Language Models and Open-Vocabulary Object Detection frameworks [10], [11]. By aligning visual features with natural language text embeddings, these open-vocabulary frameworks can recognize arbitrary categories specified via text prompts without requiring domain-specific retraining. Among recent developments, YOLO-World [12] emerges as a state-of-the-art solution, combining the real-time efficiency of standard YOLO architectures with open-vocabulary, zero-shot detection capabilities. This enables the model to identify unfamiliar objects on the fly, drastically improving situational awareness.

The primary goal of this research is to comprehensively evaluate the capacity of open-vocabulary object detection systems to mitigate the inherent safety hazards posed by closed-set limitations in autonomous driving. Specifically, the zero-shot detection reliability of the YOLO-World architecture is empirically assessed against standard YOLO models when confronted with unfamiliar or unlearned objects. To conduct this evaluation, a specific data subset derived from the diverse BDD100K driving dataset, referred to as BDD10K, is utilized.

To achieve this goal, a structured experimental methodology is employed, as illustrated in Fig. ?? . Initially, the BDD10K subset undergoes necessary dataset pre-processing, including category alias mapping, to ensure label consistency, as shown in Fig. ??a. Following this preparation, the dataset is systematically divided into known classes, such as car, bus, and truck, used exclusively for model training, and unknown classes, including pedestrian, traffic light, and bicycle, reserved to test

zero-shot capabilities. The models are then trained across multiple scales, specifically Small, Medium, and Large variants, using four distinct training configurations: two dedicated to YOLO-World training, alongside two baseline configurations for standard YOLO training, as depicted in Fig. ??b. Finally, the performance of both YOLO-World and standard YOLO architectures is rigorously measured, as detailed in Fig. ??c. This evaluation focuses on zero-shot metrics calculation, including mAP_{50} , alongside class-agnostic misclassification analysis. Through this methodology, the critical necessity of integrating vision-language features into autonomous perception systems to enhance real-time safety is quantitatively demonstrated.

II. RELATED WORK

The evolution of real-time object detection has been significantly shaped by the You Only Look Once architectures [?], [?], [3], [4]. These architectures traditionally operated under closed-vocabulary settings, requiring predefined classes and extensive retraining for novel objects. Recent advancements have introduced open-vocabulary and zero-shot detection capabilities, prominently demonstrated by YOLO-World [?]. By integrating vision-language models and prompt-based localization [?], [?], [10], YOLO-World enables real-time identification of novel objects without the need for exhaustive retraining. This shift towards open-world settings addresses critical limitations in dynamic environments where unknown classes frequently appear [?], [?], [11].

The robustness and generalization of these detection models heavily rely on the diversity and complexity of the training datasets. Autonomous driving and unconstrained environment tasks utilize comprehensive datasets such as BDD100K [?], Cityscapes [?], and IDD [?]. These datasets provide heterogeneous multi-task learning scenarios and complex urban scenes. Furthermore, datasets like ACDC [?] and CODA [?] are essential for evaluating model performance under adverse weather conditions and critical road corner cases. Utilizing these varied datasets ensures that models can handle real-world corruptions and distribution shifts effectively [?], [?].

Recent studies have aggressively adapted zero-shot and prompt-guided models to specialized domains. Mullins et al. [?] conducted an analysis of zero-shot detection capabilities to accelerate image annotation in wild blueberry cropping systems. Their study evaluated YOLO-World against Grounding DINO, revealing that Grounding DINO achieved a higher mIoU of 0.642 compared to YOLO-World at 0.503 for ripe berries. In the manufacturing sector, Tang et al. [?] explored prompt-guided AI analytics for manufacturing digital twins, successfully adapting YOLO-World to industrial object detection tasks. These implementations highlight the versatility of vision-language architectures in highly specific applications outside traditional autonomous driving contexts. However, while these domain-specific adaptations demonstrate broad flexibility, their reliance on substantial computational overhead or suboptimal zero-shot precision underscores a critical gap in developing lightweight, robust open-vocabulary models tailored for complex, unconstrained environments.

Concurrently, optimizations for object detection on edge devices continue to prioritize efficiency and environmental robustness. Gupta et al. [?] addressed the strict energy and latency constraints of edge hardware by proposing an energy-adaptive neural architecture search framework for Advanced Driver Assistance Systems. Their approach identifies highly energy-efficient base architectures, producing CNN-based models ranging from 1.8M to 2.6M parameters that achieve mean Average Precision (mAP) scores between 29.7% and 33.60% on the Cityscapes dataset, and between 20.08% and 23.35% on the BDD10K dataset. Geever et al. [?] evaluated interpolation methods for image downsampling in automotive computer vision to improve real-time inference speed. They demonstrated that while downsampling impacts performance, the differences among methods like Bilinear and Lanczos4 generally affect the average precision metric by less than 2% for most detection models, allowing for inference speed optimization with minimal accuracy trade-offs. Furthermore, Li et al. [?] proposed YOLOv7-VD, enhancing vehicle detection accuracy in complex environments by modifying the architecture to better handle difficult lighting and spatial constraints. Despite these hardware-aware optimizations successfully reducing latency, they predominantly target traditional closed-set frameworks, revealing a significant research gap in effectively deploying resource-intensive open-vocabulary vision-language models directly onto edge devices.

Despite these significant advancements, a noticeable gap remains in seamlessly unifying open-vocabulary detection with ultra-efficient edge deployment for unconstrained autonomous driving. While models like YOLO-World offer exceptional zero-shot capabilities, their performance in highly specialized, safety-critical domains often requires further architectural optimization to match domain-specific models. Additionally, current energy-aware optimizations and downsampling techniques primarily target closed-set networks. Therefore, future research must bridge this gap by developing lightweight, open-vocabulary architectures that maintain high precision on edge devices under complex, unpredictable environmental conditions.

III. RESEARCH METHODOLOGY

A. Research Architecture

The foundational difference between the comparative models in this study lies not merely in their internal feature extractors, but in their input modalities and classification paradigms, as conceptualized in Fig. ?. Standard YOLO architectures shown in Fig. ??a operate on a strictly closed-set, single-modality framework. They ingest raw visual data and map extracted features to a rigid, predefined set of classification heads established during training. Consequently, when encountering out-of-distribution objects, the model is structurally constrained to either ignore them as background or forcibly misclassify them. In contrast, YOLO-World depicted in Fig. ??b introduces an open-vocabulary, dual-modality paradigm. It simultaneously processes visual inputs alongside dynamic semantic text prompts. By leveraging vision-language

alignment, the model conditions its visual feature extraction on the provided text embeddings, allowing it to localize and classify novel objects on the fly without requiring structural retraining.

To empirically evaluate this architectural disparity within an autonomous driving context, this research employs a multi-scheme evaluation pipeline utilizing the BDD 10K dataset, detailed in Fig. ?? . The methodology systematically tests both architectures across three scaling variants: Small (s), Medium (m), and Large (l). The experimental flow is bifurcated into two primary data configurations. The first track establishes a performance upper bound by training both models on the full 10-class dataset. Following this all-class optimization, both models undergo a comprehensive evaluation to benchmark their baseline localization precision under optimal, fully-supervised conditions.

The core of the methodology, however, lies in the second training configuration, designed to isolate and test the models' responses to unseen data. In this pathway, models are trained exclusively on a restricted subset of three known classes, intentionally masking the remaining seven classes from the optimization process. During the subsequent evaluation phase, the pipelines diverge sharply where YOLO-World utilizes its zero-shot text-prompting capabilities to detect the withheld categories, while Standard YOLO undergoes a targeted misclassification and miss-rate analysis to quantify its failure modes. This structured pipeline ensures an equitable and isolated comparison between open-vocabulary adaptability and the inherent safety limitations of closed-set detection systems.

IV. PREPARE YOUR PAPER BEFORE STYLING

Before you begin to format your paper, first write and save the content as a separate text file. Complete all content and organizational editing before formatting. Please note sections IV-A to IV-H below for more information on proofreading, spelling and grammar.

Keep your text and graphic files separate until after the text has been formatted and styled. Do not number text heads— $\LaTeX$ will do that for you.

A. Abbreviations and Acronyms

Define abbreviations and acronyms the first time they are used in the text, even after they have been defined in the abstract. Abbreviations such as IEEE, SI, MKS, CGS, ac, dc, and rms do not have to be defined. Do not use abbreviations in the title or heads unless they are unavoidable.

B. Units

- Use either SI (MKS) or CGS as primary units. (SI units are encouraged.) English units may be used as secondary units (in parentheses). An exception would be the use of English units as identifiers in trade, such as “3.5-inch disk drive”.
- Avoid combining SI and CGS units, such as current in amperes and magnetic field in oersteds. This often leads to confusion because equations do not balance

dimensionally. If you must use mixed units, clearly state the units for each quantity that you use in an equation.

- Do not mix complete spellings and abbreviations of units: “Wb/m²” or “webers per square meter”, not “webers/m²”. Spell out units when they appear in text: “. . . a few henries”, not “. . . a few H”.
- Use a zero before decimal points: “0.25”, not “.25”. Use “cm³”, not “cc”.)

C. Equations

Number equations consecutively. To make your equations more compact, you may use the solidus (/), the exp function, or appropriate exponents. Italicize Roman symbols for quantities and variables, but not Greek symbols. Use a long dash rather than a hyphen for a minus sign. Punctuate equations with commas or periods when they are part of a sentence, as in:

$$a + b = \gamma \tag{1}$$

Be sure that the symbols in your equation have been defined before or immediately following the equation. Use “(1)”, not “Eq. (1)” or “equation (1)”, except at the beginning of a sentence: “Equation (1) is . . .”

D. $\LaTeX$ -Specific Advice

Please use “soft” (e.g., `\eqref{Eq}`) cross references instead of “hard” references (e.g., (1)). That will make it possible to combine sections, add equations, or change the order of figures or citations without having to go through the file line by line.

Please don't use the `{eqnarray}` equation environment. Use `{align}` or `{IEEEeqnarray}` instead. The `{eqnarray}` environment leaves unsightly spaces around relation symbols.

Please note that the `{subequations}` environment in $\LaTeX$ will increment the main equation counter even when there are no equation numbers displayed. If you forget that, you might write an article in which the equation numbers skip from (17) to (20), causing the copy editors to wonder if you've discovered a new method of counting.

$\BibTeX$ does not work by magic. It doesn't get the bibliographic data from thin air but from .bib files. If you use $\BibTeX$ to produce a bibliography you must send the .bib files.

$\LaTeX$ can't read your mind. If you assign the same label to a subsection and a table, you might find that Table I has been cross referenced as Table IV-B3.

$\LaTeX$ does not have precognitive abilities. If you put a `\label` command before the command that updates the counter it's supposed to be using, the label will pick up the last counter to be cross referenced instead. In particular, a `\label` command should not go before the caption of a figure or a table.

Do not use `\nonumber` inside the `{array}` environment. It will not stop equation numbers inside `{array}` (there won't be any anyway) and it might stop a wanted equation number in the surrounding equation.

E. Some Common Mistakes

- The word “data” is plural, not singular.
- The subscript for the permeability of vacuum μ_0 , and other common scientific constants, is zero with subscript formatting, not a lowercase letter “o”.
- In American English, commas, semicolons, periods, question and exclamation marks are located within quotation marks only when a complete thought or name is cited, such as a title or full quotation. When quotation marks are used, instead of a bold or italic typeface, to highlight a word or phrase, punctuation should appear outside of the quotation marks. A parenthetical phrase or statement at the end of a sentence is punctuated outside of the closing parenthesis (like this). (A parenthetical sentence is punctuated within the parentheses.)
- A graph within a graph is an “inset”, not an “insert”. The word alternatively is preferred to the word “alternately” (unless you really mean something that alternates).
- Do not use the word “essentially” to mean “approximately” or “effectively”.
- In your paper title, if the words “that uses” can accurately replace the word “using”, capitalize the “u”; if not, keep using lower-cased.
- Be aware of the different meanings of the homophones “affect” and “effect”, “complement” and “compliment”, “discreet” and “discrete”, “principal” and “principle”.
- Do not confuse “imply” and “infer”.
- The prefix “non” is not a word; it should be joined to the word it modifies, usually without a hyphen.
- There is no period after the “et” in the Latin abbreviation “et al.”.
- The abbreviation “i.e.” means “that is”, and the abbreviation “e.g.” means “for example”.

An excellent style manual for science writers is [?].

F. Authors and Affiliations

The class file is designed for, but not limited to, six authors. A minimum of one author is required for all conference articles. Author names should be listed starting from left to right and then moving down to the next line. This is the author sequence that will be used in future citations and by indexing services. Names should not be listed in columns nor group by affiliation. Please keep your affiliations as succinct as possible (for example, do not differentiate among departments of the same organization).

G. Identify the Headings

Headings, or heads, are organizational devices that guide the reader through your paper. There are two types: component heads and text heads.

Component heads identify the different components of your paper and are not topically subordinate to each other. Examples include Acknowledgments and References and, for these, the correct style to use is “Heading 5”. Use “figure caption” for your Figure captions, and “table head” for your table title. Run-in heads, such as “Abstract”, will require you

to apply a style (in this case, italic) in addition to the style provided by the drop down menu to differentiate the head from the text.

Text heads organize the topics on a relational, hierarchical basis. For example, the paper title is the primary text head because all subsequent material relates and elaborates on this one topic. If there are two or more sub-topics, the next level head (uppercase Roman numerals) should be used and, conversely, if there are not at least two sub-topics, then no subheads should be introduced.

H. Figures and Tables

a) *Positioning Figures and Tables:* Place figures and tables at the top and bottom of columns. Avoid placing them in the middle of columns. Large figures and tables may span across both columns. Figure captions should be below the figures; table heads should appear above the tables. Insert figures and tables after they are cited in the text. Use the abbreviation “Fig. ??”, even at the beginning of a sentence.

TABLE I
TABLE TYPE STYLES

Table Head	Table Column Head		
	Table column subhead	Subhead	Subhead
copy	More table copy ^a		

^aSample of a Table footnote.

Figure Labels: Use 8 point Times New Roman for Figure labels. Use words rather than symbols or abbreviations when writing Figure axis labels to avoid confusing the reader. As an example, write the quantity “Magnetization”, or “Magnetization, M”, not just “M”. If including units in the label, present them within parentheses. Do not label axes only with units. In the example, write “Magnetization (A/m)” or “Magnetization {A[m(1)]}”, not just “A/m”. Do not label axes with a ratio of quantities and units. For example, write “Temperature (K)”, not “Temperature/K”.

ACKNOWLEDGMENT

The preferred spelling of the word “acknowledgment” in America is without an “e” after the “g”. Avoid the stilted expression “one of us (R. B. G.) thanks ...”. Instead, try “R. B. G. thanks...”. Put sponsor acknowledgments in the unnumbered footnote on the first page.

REFERENCES

- [1] H.-H. Jebamikyous and R. F. Kashef, “Autonomous vehicles perception (avp) using deep learning: Modeling, assessment, challenges,” *IEEE Access*, vol. PP, pp. 1–1, 2022.
- [2] A. Pravallika, M. Hashmi, and A. Gupta, “Deep learning frontiers in 3d object detection: A comprehensive review for autonomous driving,” *IEEE Access*, vol. 12, pp. 173 936–173 980, 2024.
- [3] J. Redmon, S. Divvala, R. B. Girshick, and A. Farhadi, “You only look once: Unified, real-time object detection,” *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 779–788, 2015.
- [4] J. R. Terven, D.-M. Córdova-Esparza, and J.-A. Romero-González, “A comprehensive review of yolo architectures in computer vision: From yolov1 to yolov8 and yolo-nas,” *Mach. Learn. Knowl. Extr.*, vol. 5, pp. 1680–1716, 2023.

- [5] F. Yu, H. Chen, X. Wang, W. Xian, Y. Chen, F. Liu, V. Madhavan, and T. Darrell, "Bdd100k: A diverse driving dataset for heterogeneous multitask learning," *2020 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2633–2642, 2018.
- [6] A. Dhamija, M. Günther, J. Ventura, and T. Boulton, "The overlooked elephant of object detection: Open set," in *2020 IEEE Winter Conference on Applications of Computer Vision (WACV)*, 2020, pp. 1010–1019.
- [7] Z. Li, X. Kang, J. West, and K. Khoshelham, "Out-of-distribution detection in 3d applications: a review," *ArXiv*, vol. abs/2507.00570, 2025.
- [8] K. Li, K. Chen, L. Hong, C. Ye, J. Han, Y. Chen, W. Zhang, C. Xu, D.-Y. Yeung, X. Liang, Z. Li, and H. Xu, "Coda: A real-world road corner case dataset for object detection in autonomous driving," in *ArXiv*, vol. abs/2203.07724, 2022.
- [9] D. Chen, Z. Zhang, Y. Liu, and X. Yang, "Insight: Enhancing autonomous driving safety through vision-language models on context-aware hazard detection and edge case evaluation," *ArXiv*, vol. abs/2502.00262, 2025.
- [10] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *ArXiv*, vol. abs/2103.00020, 2021.
- [11] Z. Li, Z. Xiang, J. West, and K. Khoshelham, "From open vocabulary to open world: Teaching vision language models to detect novel objects," *ArXiv*, vol. abs/2411.18207, 2024.
- [12] T. Cheng, L. Song, Y. Ge, W. Liu, X. Wang, and Y. Shan, "Yolo-world: Real-time open-vocabulary object detection," *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 16 901–16 911, 2024.